

SAM 3D Animal: Promptable Animal 3D Reconstruction from Images in the Wild

Xuyi Hu^{1*} Jin Lyu^{2*} Jiuming Liu¹ Yebin Liu³ Silvia Zuffi⁴ Liang An^{3†} Stefan Goetz¹

¹University of Cambridge ²Southern University of Science and Technology ³Tsinghua University ⁴IMATI-CNR, Milan, Italy



Figure 1: A promptable view of multi-animal 3D reconstruction. We present **SAM 3D Animal**, a promptable framework that addresses the problem of joint 3D reconstruction of multiple animals from a single image. Here, we show the input image together with the overlay reconstruction results.

Abstract

3D animal reconstruction in the wild remains challenging due to large species variation, frequent occlusions, and the prevalence of multi-animal scenes, while existing methods predominantly focus on single-animal settings. We present SAM 3D Animal, the first promptable framework for multi-animal 3D reconstruction from a single image. Built on the SMAL+ parametric animal model, our method jointly reconstructs multiple instances and supports flexible prompts in the form of keypoints and masks which enable more reliable disambiguation in crowded and occluded scenes. To train such a model, we further introduce Herd3D, a multi-animal 3D dataset containing over 5K images, designed to increase diversity in species, interactions, and occlusion patterns. Experiments on the Animal3D, APTv2, and Animal Kingdom datasets show that our framework achieves state-of-the-art results over both existing model-based and model-free methods, demonstrating a scalable and effective solution for prompt-driven animal 3D reconstruction in the wild.

1 Introduction

Animals are a fundamental part of the visual world, yet 3D reconstruction research remains heavily skewed toward humans. Human-centric methods have advanced pose and shape estimation dramatically [18, 49, 48, 15, 36, 4, 23]. In contrast, animal reconstruction still suffers from scarce datasets, broad species variation, and inconsistent anatomical definitions.

Parametric animal models such as SMAL [52] and SMAL+ [50] provide an effective basis for recovering 3D pose and shape from a single image [51, 41, 30, 28, 1]. These approaches typically

* Equal contribution.

† Corresponding author.

focus on one animal at a time and often rely on pre-cropped inputs or strong object detections. However, many in-the-wild animal scenes contain multiple individuals with mutual occlusion, and complex interactions that invalidate single-animal assumptions.

Multi-animal 3D reconstruction raises unique challenges beyond those of the single-object case. First, instance association becomes ambiguous when animals overlap or occlude one another. Second, pose and shape estimation must be jointly consistent across multiple hypotheses, since mistakes on one individual can be amplified by false depth ordering or incorrect occlusion reasoning. Third, available datasets rarely provide dense multi-animal 3D annotations, which hinders supervised learning for crowded scenes.

To overcome these challenges, we draw inspiration from promptable reconstruction in human vision. Recent works such as SAM 3D Body [44] demonstrates that explicit prompts can guide a model to focus on a desired subject and resolve ambiguity in cluttered scenes. Prompts can take the form of keypoints or masks, each providing a different level of spatial and semantic guidance.

In this paper, we introduce **SAM 3D Animal**, the first promptable framework for multi-animal 3D reconstruction, see Fig. 1. Our model uses the SMAL+ template [50] and can ingest optional prompts in two modalities: keypoints for skeletal alignment and masks for precise silhouette discrimination. This promptable design allows SAM 3D Animal to recover multiple animals jointly from a single image. Different from SAM 3D Body, which reconstructs a single prompted subject per forward pass, our model adopts a set-prediction paradigm that recovers all animal instances in one shot via DETR-style [10] bipartite matching, eliminating the need for per-instance bounding-box cropping.

However, training such a multi-instance model with only 2D-annotated datasets is insufficient, as 2D keypoints and silhouettes alone cannot provide the per-instance 3D shape supervision needed to resolve inter-animal occlusions. To address this, we propose Herd3D, a multi-animal 3D dataset containing over 5K images with per-instance ground-truth meshes, designed to increase diversity in species, interactions, and occlusion patterns. The generation pipeline of Herd3D is adapted from GenZoo [30], thus each animal is naturally labeled with image-aligned SMAL+ model.

To demonstrate the effectiveness of SAM 3D Animal, we compare it with state-of-the-art animal mesh recovery methods on publicly available datasets including Animal3D [41], APTv2 [45], and Animal Kingdom [29]. Even without any prompt, our model achieves competitive or superior results compared to existing methods. When prompts are provided, performance improves consistently across all benchmarks, with up to 54% AP gain and 80% mAP gain on the out-of-domain Animal Kingdom dataset over the strongest baseline, as well as 5.2 PA-MPJPE improvement on Animal3D. Ablation studies confirm that Herd3D brings consistent improvements, particularly on multi-animal benchmarks, and that keypoint prompts are the dominant contributor among prompt modalities, with performance scaling monotonically as the number of keypoints increases.

2 Related Work

2.1 Model-Free Reconstruction

Model-free animal reconstruction learns 3D structure directly from image or video collections without assuming a predefined template [14, 40]. Early methods model category-specific articulated animals from single-view image collections by separating a predefined skeleton prior from instance-specific deformations [47, 39, 38, 16]. Later approaches extend to a wider variety of species, either by learning a unified shape model [24] or by applying linear skinning to deform learned 3D object shapes [2]. However, these methods still struggle with extreme poses, heavy occlusions, and limited viewpoint coverage, often producing geometrically ambiguous reconstructions.

2.2 Model-Based Reconstruction

Model-based animal reconstruction typically relies on predefined quadruped templates such as SMAL [52] that encode the shape and articulation structure of specific animal families. These models either fit predefined 3D templates to animal images using 2D observations such as keypoints or silhouettes [51, 6, 7], or directly reconstruct the 3D shape from image or video observations [11, 43, 47, 53, 27]. This parametric formulation offers interpretable and controllable representations, which make the reconstructed animals readily animatable and editable. Recent works further extend

SMAL-based reconstruction to broader quadruped species and training settings. AWOL [50] maps CLIP-style embeddings to the SMAL⁺ parameter space for language- and image-guided animal shape generation. RAW [20] reconstructs animals jointly with their surrounding environment, including multi-animal scenes. However, it relies on rigid animal assets rather than articulated animal models, and therefore does not address fine-grained articulated animal reconstruction.

2.3 Animal Pose Estimation Datasets

In comparison to humans, the construction of large-scale animal datasets is significantly more challenging because animals are difficult to capture in controlled environments and exhibit substantial morphological diversity across species. Animal benchmarks, such as Stanford Extra [5], Animal Pose [8], and AwA-Pose [3], remain limited to 2D annotations. Existing 3D animal datasets, such as Animal3D [41], CtrlAni3D [28], GenZoo [30] and FemaleSaanenGoat [17], predominantly focus on single-animal instances, whereas large-scale benchmarks like APT-36K [46], and Animal Kingdom [29] provide only 2D annotations. This limitation restricts the development of methods that can jointly model inter-animal occlusions, spatial relationships, and pose dependencies in multi-animal scenes.

2.4 Promptable Mesh Reconstruction

Promptable mesh reconstruction has recently emerged in human mesh recovery, where auxiliary cues guide 3D estimation under occlusion and crowding. PromptHMR [35] incorporates full-image context with spatial and semantic prompts for pose and shape estimation. SAM 3D Body [44] extends this idea to full-body recovery through a promptable encoder-decoder architecture supporting keypoint and mask prompts. SAM-Body4D [13] further leverages temporally consistent masklets to produce coherent mesh trajectories from videos. However, these methods are designed for humans, whereas our work targets the animal domain, where large morphological diversity, inter-animal occlusion, and multi-instance interactions must be jointly considered.

3 SAM 3D Animal Model

3.1 Preliminary

SMAL+. SMAL+ [50], denoted as $\mathcal{M}(\beta, \theta, \gamma)$, extends the original SMAL [52] model by incorporating training samples from D-SMAL [32] and hSMAL [21], alongside new species such as the giraffe, bear, mouse and rat. This results in a broader, 145-dimensional shape space learned from a total of 145 animals. The inputs to SMAL+ are the shape parameters $\beta \in \mathbb{R}^{145}$ and the pose parameters $\theta \in \mathbb{R}^{35 \times 3}$ (using an axis-angle representation) and the global translation $\gamma \in \mathbb{R}^3$. By applying linear blendshapes and Linear Blend Skinning (LBS), SMAL+ outputs a posed mesh with vertices $\mathbf{V} \in \mathbb{R}^{3889 \times 3}$ and faces $\mathbf{F} \in \mathbb{N}^{7774 \times 3}$.

3.2 End-to-end Multi-Instance Network

Given an animal image, our model can reconstruct all animals in the image without requiring preprocessed bounding boxes, and support masks or keypoints as prompts.

Encoder. Starting from the image $x \in \mathbb{R}^{H \times W \times 3}$, we utilize the ViT-Huge Encoder [12] to generate the feature tokens $\mathcal{F} \in \mathbb{R}^{(H_0 \times W_0) \times C_0}$, where C_0 , H_0 , and W_0 are the channels, the height and width of the feature map, respectively. In our case, $W = H = 512$, $C_0 = 1280$, $W_0 = H_0 = 32$.

Decoder. Inspired by SAM 3D Body [44], the decoder is a SAM-style promptable Transformer (see Fig. 2). Specifically, it takes the feature tokens $\mathcal{F}$ and a set of query tokens as input, then performs cross-attention, and finally predicts the SMAL+ parameters, cameras, and bounding boxes. Note that, different from SAM 3D Body which predicts single person at a time, our model directly predict $P = 30$ possible instances at a time, eliminating the need for bounding box input. The query tokens consist of six distinct token groups for decoder layer l ($0 < l \leq 6$):

$$\mathcal{Q}^l = \text{Concat}(\mathcal{Q}_{\text{params}}^l, \mathcal{Q}_{\text{box}}^l, \mathcal{Q}_{2D}^l, \mathcal{Q}_{3D}^l, \mathcal{Q}_{\text{prompt}}^l) \in \mathbb{R}^{N \times D}, \quad (1)$$




Figure 2: SAM 3D Animal Model structure.

where Q_{params}^l , Q_{box}^l , Q_{2D}^l , Q_{3D}^l and Q_{prompt}^l represent the initial SMAL+ pose tokens, bounding box tokens, 2D keypoints tokens, 3D keypoints tokens, the interaction prompt tokens. Note that feature dimension $D = 1024$. $N = P \times 405 = 12150$ where 405 is full token dimension for each prediction.

During the forward pass, query tokens interact with the flattened image features $\mathcal{F}$ through a standard multi-head cross-attention mechanism. At layer l , we first concatenate Q^l with its previous state $Q^{(l-1)}$ to get Q^c , and the attention operation is defined as:

$$Q^{(l+1)} = \text{CrossAttention}(Q^c, \mathcal{F}) = \text{Softmax}\left(\frac{(Q^c W_Q)(W_K \mathcal{F})}{\sqrt{d_k}} (F W_V)\right), \quad (2)$$

where W_Q , W_K , and W_V are the learnable projection matrices for the queries, keys, and values, and d_k is the scaling factor based on the head dimension. At first layer, Q^0 is randomly initialized.

A critical feature of this architecture is the layer-wise keypoint feedback loop. After cross attention, the model further explicitly refreshes the 2D and 3D keypoint tokens for the subsequent layer $(l+1)$ using the current predictions. For 2D keypoints, the tokens are augmented using both positional embeddings of the predicted coordinates and local image features sampled at those locations:

$$Q_{2D}^{(l+1)} \leftarrow Q_{2D}^{(l+1)} + \phi_{\text{pos}}(x_{2D}^{(l+1)}) + \phi_{\text{feat}}(\mathcal{F}(x_{2D}^{(l+1)})), \quad (3)$$

where ϕ_{pos} and ϕ_{feat} denote linear projections, and $\mathcal{F}(x_{2D}^l)$ represents the image features sampled at the predicted 2D locations. In parallel, the 3D keypoint tokens are updated purely based on the geometric embeddings of the normalized 3D coordinates:

$$Q_{3D}^{(l+1)} \leftarrow Q_{3D}^{(l+1)} + \psi_{\text{pos}}(x_{3D}^{(l+1)}), \quad (4)$$

where ψ_{pos} is the linear projection mapping the 3D coordinates into the token embedding space. This iterative mechanism ensures that subsequent layers are conditioned on the most recent geometric and appearance estimates, facilitating the precise convergence of the final output meshes and keypoint projections. It is worth mentioning that only Q_{params} and Q_{box} at the final layer are used for generating predictions. “params” in Q_{params} and Fig. 2 refers to both SMAL+ parameters and camera parameters.

Bipartite Matching for Multi-Animal Instances. To enable end-to-end training without heuristic post-processing such as Non-Maximum Suppression (NMS), we adopt a set prediction formulation following the DETR paradigm [10]. Specifically, we employ bipartite matching via the Hungarian algorithm [19] to find the optimal one-to-one assignment between the fixed-size set of P predicted animal hypotheses and the M ground-truth instances. The matching cost is a weighted combination of bounding box $\mathcal{L}_1$ distance, Generalized IoU [31], focal-style confidence penalty [33], and masked 2D keypoint distance. Once the optimal assignment is established, predicted outputs are reordered for loss computation. See Appendix for more details.

3.3 Loss Functions

After establishing the correspondence between predicted outputs and ground-truth labels via bipartite matching, we optimize the model using a multi-task loss function formulated as:

$$\mathcal{L} = \lambda_{\text{params}} \mathcal{L}_{\text{params}} + \lambda_{2D} \mathcal{L}_{2D} + \lambda_{3D} \mathcal{L}_{3D} + \lambda_{\text{box}} \mathcal{L}_{\text{box}}, \quad (5)$$

where $\lambda_{\{\cdot\}}$ denotes the weighting coefficients used to balance the respective loss contributions.

Parameter Loss ($\mathcal{L}_{\text{params}}$) computes the ℓ_2 distance between the predicted SMAL+ shape and pose parameters and their corresponding ground-truth values if provided.

Keypoint Losses ($\mathcal{L}_{2\text{D}}$, $\mathcal{L}_{3\text{D}}$) represent the ℓ_1 distance between the ground-truth 2D and 3D keypoint positions and the ones regressed from predicted SMAL+, respectively.

Bounding Box Loss ($\mathcal{L}_{\text{box}}$) supervises localization accuracy through a combination of coordinate regression, geometric alignment, confidence scoring and the denoising training strategy of DN-DETR [22]:

$$\mathcal{L}_{\text{box}} = \mathcal{L}_{\text{coord}} + \mathcal{L}_{\text{giou}} + \mathcal{L}_{\text{conf}} + \mathcal{L}_{\text{dn}}, \quad (6)$$

where $\mathcal{L}_{\text{coord}}$ is an ℓ_1 loss over normalized bounding box coordinates, and $\mathcal{L}_{\text{giou}}$ refers to the Generalized IoU (GIoU) loss [31]. To refine the objectness score, $\mathcal{L}_{\text{conf}}$ employs a Binary Cross-Entropy (BCE) loss, where the actual IoU between the matched predicted and ground-truth boxes serves as the soft target for the predicted confidence.

4 Herd3D Dataset




Figure 3: Example from the Herd3D dataset. This figure shows a generated scene with eight dogs, together with the corresponding canny map, depth map, RGB image, and 3D reconstruction overlay. The full Herd3D dataset contains multi-animal scenes with 2 to 8 animals per image.

To support multi-animal 3D reconstruction in real-world scenarios, we construct Herd3D, a large-scale dataset specifically designed for multi-animal scenes which contains over 5K images and 118 species (see Fig. 3). We believe that GenZoo [30] provides a strong and practical starting point for constructing large-scale animal datasets, because it couples a parametric animal model with controllable image synthesis, which enables scalable generation of paired images and geometry while maintaining pose and shape consistency. Building on GenZoo, we adapt the pipeline for multi-animal data generation. To construct group layouts, we sample up to 8 animals per image and place them on a shared ground plane. For each instance, we set $t_y = 0$ and sample translations with $t_x \in [-1.5, 1.5]$ using non-adjacent horizontal bins to limit excessive overlap, and $t_z \in [8, 50]$ from predefined depth intervals while constraining the group depth span to at most 30; we add a small x and z jitter within ± 1.5 and apply a fixed ground alignment offset of 0.3. We further diversify global orientation by sampling pitch in $[-15^\circ, 15^\circ]$ and yaw in $[0^\circ, 360^\circ]$. To accommodate the increased complexity of multi-animal scenes-including frequent occlusions, higher ambiguity in instance-wise orientation and limited pose diversity, we adapt the GenZoo pipeline with several targeted modifications. We (i) impose scene layout constraints by placing all animals on a shared ground plane, (ii) expand the pose pool by integrating Animal3D poses [41] to increase pose diversity, (iii) replace the ControlNet backend with Qwen-Image-ControlNet-Union [37] to better preserve geometry and occlusion ordering, and (iv) resolve multi-animal orientation ambiguity via a two-stage Qwen3-VL-8B-Instruct [34] prompting scheme, which first predicts left-to-right per-animal facing directions and then composes a single coherent final prompt that integrates the species information, camera settings, and scene attributes. Each synthetic image has a resolution of 1024×1024 and includes annotations for SMAL+ parameters, 2D keypoints, 3D keypoints and bounding boxes.

5 Experiments

Datasets. We curate a comprehensive training corpus of 49.2K images containing both 2D and 3D annotations. Specifically, we aggregate the training splits of Animal Pose [8], APTv2 [45], AwA-Pose [3], Stanford Extra [5], Animal3D [41], and our newly introduced Herd3D. For evaluation, following the protocol in AniMer [28], we report results on two in-domain datasets (Animal3D and APTv2) alongside an out-of-domain (OOD) dataset Animal Kingdom.



Figure 4: Qualitative comparisons on Animal3D, Animal Kingdom and APT-36K datasets. We compare our results with AniMer [28], GenZoo [30] and 3D Fauna [24]. AniMer, GenZoo, and 3D Fauna require pre-cropped single-animal images as input, whereas our method directly processes the full image and retrieves reconstructions whose confidence exceeds a predefined threshold, eliminating the need for cropping.

Baselines. We benchmark our approach against three recent state-of-the-art (SOTA) methods to ensure a comprehensive evaluation across different architectural paradigms. For model-based techniques, we compare with AniMer [28], a transformer-based architecture utilizing the SMAL model, and GenZoo [30], which builds upon the SMAL+ variant. Additionally, we include 3D Fauna [24] as a representative SOTA model-free reconstruction approach.

Evaluation Metrics. We evaluate 3D accuracy using the Procrustes-Aligned Mean Per Joint Position Error (PA-MPJPE). For 2D accuracy, we report the Percentage of Correct Keypoints (PCK), AP (Average Precision) and mAP (mean Average Precision) [25].

Implementation Details. Our network is optimized using AdamW [26] with an initial learning rate of 2×10^{-5} , incorporating a linear warmup over the first 15 epochs. Similar to AniMer [28], we employ a two-stage training strategy consisting of 250 epochs for the first stage and 250 epochs for the second. We apply prompt dropout for robustness: the mask prompt is dropped with 50% probability, the entire keypoint prompt is dropped with a probability of $p_{\text{drop}} = 0.2$, and otherwise each keypoint is independently masked out with a rate sampled uniformly from $[0, 0.7]$ per step, encouraging the model to handle partial or absent prompts at inference time. Training is distributed across four RTX 4090 GPUs with a gradient accumulation step of 16. To balance the objective function, our empirical loss weighting factors are set to $\lambda_{\text{params}} = 1$, $\lambda_{2D} = 5$, $\lambda_{3D} = 5$, and $\lambda_{\text{box}} = 1$.

5.1 Comparison

Comparison without prompts. We present the quantitative results in Table 1. Without any prompt, our method already achieves competitive or superior performance relative to existing approaches. On Animal3D, the prompt-free variant attains a PA-MPJPE of 80.7 mm, slightly outperforming AniMer (81.0 mm), while achieving a higher mAP (49.3 vs. 47.2). On APTv2, keypoint localization improves substantially, with PCK@0.1 reaching 87.9, far surpassing GenZoo (64.1) and AniMer (62.4), though AP remains lower (49.4 vs. 55.5 for GenZoo), indicating that the two paradigms exhibit complementary strengths under different metrics. On the OOD Animal Kingdom benchmark, our prompt-free results lead all metrics, demonstrating stronger generalization to unseen scenes.

Prompt-driven performance gains. A key advantage of our framework is its ability to leverage auxiliary prompts at inference time. When supplied with keypoints from an off-the-shelf ViTPose [42]



Figure 5: Qualitative evaluation of SAM 3D Animal. For each example, we show: (a) the input image, and (b) the 3D reconstruction overlay. To demonstrate the robustness of SAM 3D Animal, we visualize results across diverse animal species and challenging in-the-wild scenarios, including unusual poses, large viewpoint variations, and crowded scenes with several animals.

Table 1: **Quantitative comparisons on the Animal3D dataset, APTv2 dataset and Animal Kingdom dataset.**

Method	Animal3D			APTv2			Animal Kingdom		
	PA-MPJPE↓	AP↑	mAP↑	PCK@0.1↑	AP↑	mAP↑	PCK@0.1↑	AP↑	mAP↑
3D Fauna [24]	202.8	63.0	10.0	32.6	39.5	7.0	31.4	13.5	2.9
GenZoo [30]	123.2	85.3	34.0	64.1	55.5	19.2	42.1	23.6	4.7
AniMer [28]	81.0	93.4	47.2	62.4	51.8	20.1	58.2	39.7	10.4
Ours	80.7	92.8	49.3	87.9	49.4	23.5	59.2	45.0	12.6
Ours (w/ vitpose)	80.7	92.8	50.0	88.8	55.5	27.9	63.0	50.5	14.5
Ours (w/ GT prompt)	75.8	93.8	53.4	89.0	57.4	29.7	66.2	60.1	17.7

estimator, performance improves consistently: on APTv2, AP rises from 49.4 to 55.5 and mAP from 23.5 to 27.9; on Animal Kingdom, AP increases from 45.0 to 50.5. This practical variant already matches or surpasses the best baseline on most metrics. With ground-truth keypoint prompts, the gains become substantially larger: on APTv2, PCK@0.1 reaches 89.0 (vs. 62.4 for AniMer) and AP reaches 57.4 (vs. 55.5 for GenZoo); on Animal Kingdom, AP improves to 60.1 and mAP to 17.7, roughly doubling AniMer’s 10.4. These results confirm that prompting provides a scalable mechanism for improving reconstruction quality, with performance increasing monotonically as prompt fidelity improves—a unique advantage that existing methods cannot replicate.

Qualitative comparison. Fig. 4 presents visual comparisons across the three benchmarks. 3D Fauna, as a model-free approach, produces coarse reconstructions that lack geometric detail. GenZoo and AniMer yield plausible shapes but exhibit less accurate alignment with the input image. Our method consistently produces reconstructions that are better aligned with the observed pose and viewpoint. Additional results spanning diverse species and challenging in-the-wild scenarios are shown in Fig. 5.




Figure 6: Ablation studies. Keypoint prompting, mask prompting, and training with our Herd3D dataset each lead to improved performance, as discussed in Sec. 5.2.

5.2 Ablation

Table 2: Ablation study on the Animal3D dataset, APTv2 dataset and Animal Kingdom dataset.

Variant	Animal3D			APTv2			Animal Kingdom		
	PA-MPJPE↓	AP↑	mAP↑	PCK@0.1↑	AP↑	mAP↑	PCK@0.1↑	AP↑	mAP↑
w/o Herd3D	79.5	<u>93.8</u>	51.6	88.0	48.7	24.7	62.5	51.1	15.7
w/o prompt	80.7	92.8	49.3	87.9	49.4	23.5	59.2	45.0	12.6
w/o kp	80.8	93.1	49.8	87.9	49.7	23.5	59.3	45.4	12.8
w/o mask	75.8	94.1	53.4	88.9	57.3	29.6	<u>66.0</u>	60.1	17.5
Full	75.8	<u>93.8</u>	53.4	89.0	57.4	29.7	66.2	60.1	17.7

Table 3: Ablation study on the number of prompt keypoints.

Number	Animal3D			APTv2			Animal Kingdom		
	PA-MPJPE↓	AP↑	mAP↑	PCK@0.1↑	AP↑	mAP↑	PCK@0.1↑	AP↑	mAP↑
0	80.8	93.1	49.8	87.9	49.7	23.5	59.3	45.4	12.8
5	79.5	94.1	51.7	88.9	54.6	26.7	63.9	56.9	16.3
10	78.6	94.1	51.7	89.1	56.8	28.7	65.3	58.6	17.1
15	77.4	<u>93.8</u>	52.5	<u>89.0</u>	57.5	29.6	66.0	60.2	17.7

We ablate three design axes to isolate their respective contributions: training data, prompt modality, and prompt density. Results are reported in Tables 2 and 3, with qualitative examples in Fig. 6.

Effect of Herd3D. Removing Herd3D from the training set leads to a consistent performance drop across all three benchmarks, with the largest degradation observed on APTv2 (Table 2). This is expected: Herd3D is the primary source of multi-animal scenes with per-instance annotations, and its absence disproportionately affects benchmarks that feature crowded or overlapping subjects. The result validates that our curated dataset fills a genuine gap in existing training resources rather than merely increasing data volume.

Keypoint vs. mask prompts. Comparing the prompt modality variants in Table 2 reveals a clear asymmetry: keypoint prompts are the dominant contributor, while mask prompts provide only marginal gains. Specifically, removing the mask prompt (*w/o mask*) leaves performance nearly unchanged from the full model, whereas removing keypoint prompts (*w/o kp*) reduces results to the level of the unprompted baseline. We attribute this to two factors. First, keypoints encode the articulated skeletal structure directly, which is precisely what the SMAL model needs to resolve pose ambiguity; masks, by contrast, convey only silhouette-level information that is largely redundant with the image features already extracted by the backbone. Second, the mask prompts are generated by SAM [9] at inference time, and segmentation errors, particularly on thin limbs or under occlusion, introducing noise that dilutes the prompt signal. Nonetheless, the full model still edges ahead of the keypoint-only variant on APTv2 and Animal Kingdom, suggesting that mask prompts offer a small but consistent complementary benefit when segmentation quality is adequate.

Number of keypoints. Table 3 displays the performance as a function of the number of keypoint prompts, ranging from 0 to 15. For each sample, we randomly select the specified number of keypoints from all available annotations as the prompt input. Performance improves monotonically, with the steepest gain occurring between 0 and 5 keypoints, even a sparse set of randomly chosen




Figure 7: **Performance under different visibility levels.** We group test samples by the number of visible keypoints into Low, Mid, and High buckets. (a) mAP on APTv2. (b) mAP on Animal Kingdom. Error bars denote standard deviation across visibility counts within each group.

landmarks is sufficient to substantially disambiguate pose. Beyond 5, improvements continue at a diminishing rate, indicating that the initial keypoints resolve the most salient ambiguities while additional ones refine secondary joints. This graceful scaling is practically appealing: users can supply as few or as many keypoints as available, and the model extracts value from each additional annotation without requiring a fixed-size input.

5.3 Impact of Prompt Strategy Under Occlusion

To understand how different prompt strategies behave under varying levels of occlusion, we partition the test set by the number of visible keypoints into three groups (Low, Mid, High) and evaluate three prompt modes: GT Prompt (ground-truth 2D keypoints), ViTPose Prompt (automatically detected by ViTPose [42]), and No Prompt. Results are reported using mAP on both APTv2 and Animal Kingdom, as shown in Fig. 7. Note that for Animal Kingdom, the lowest visibility group starts at 8–11 keypoints because the dataset’s annotations contain a minimum of 8 visible keypoints per instance.

Prompt matters more under heavy occlusion. Across both datasets, the relative gain of prompting over the No Prompt baseline is largest in the Low visibility group and diminishes as more keypoints become visible. On APTv2, ViTPose Prompt improves mAP over No Prompt by 57% relative in the Low group, 40% in Mid, and 26% in High. On Animal Kingdom, the same trend holds: 28% relative gain in Low, 16% in Mid, and 7% in High. This confirms that the model relies more heavily on prompt-provided spatial priors when visual evidence is limited, and that prompting offers the greatest practical benefit precisely in the most challenging scenarios.

ViTPose Prompt is a practical alternative to GT. On APTv2, ViTPose Prompt consistently performs close to the GT upper bound across all visibility levels, demonstrating that an off-the-shelf keypoint detector can serve as an effective substitute for manual annotations at inference time.

6 Conclusion

We presented SAM 3D Animal, a promptable framework for multi-animal 3D reconstruction from a single image. Unlike prior animal reconstruction methods that predominantly focus on animal-centric images, our approach reconstructs multiple animals jointly through a set-prediction formulation and supports flexible keypoint and mask prompts to resolve ambiguity in crowded and occluded scenes.

Limitation. While SAM 3D Animal shows strong performance, it remains limited by the SMAL+ shape space and is therefore mainly applicable to quadruped-like animals. Moreover, relative depth ordering between animals is not explicitly constrained, which can cause inaccurate spatial arrangements under severe occlusion. Future work could explore more flexible animal representations and explicit depth-aware scene reasoning.

References

- [1] Liang An, Jin Lyu, Li Lin, Pujin Cheng, Yebin Liu, and Xiaoying Tang. Animer+: Unified pose and shape estimation across mammalia and aves via family-aware transformer. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 48(3):3233–3249, 2026.
- [2] Mehmet Aygun and Oisín Mac Aodha. Saor: Single-view articulated object reconstruction. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 10382–10391, 2024.
- [3] Prianka Banik, Lin Li, and Xishuang Dong. A novel dataset for keypoint detection of quadruped animals from images, 2021.
- [4] Fabien Baradel, Matthieu Armando, Salma Galaaoui, Romain Brégier, Philippe Weinzaepfel, Grégory Rogez, and Thomas Lucas. Multi-hmr: Multi-person whole-body human mesh recovery in a single shot. In *European Conference on Computer Vision*, pages 202–218. Springer, 2024.
- [5] Benjamin Biggs, Oliver Boyne, James Charles, Andrew Fitzgibbon, and Roberto Cipolla. Who left the dogs out? 3d animal reconstruction with expectation maximization in the loop. In *European Conference on Computer Vision*, pages 195–211. Springer, 2020.
- [6] Benjamin Biggs, Thomas Roddick, Andrew Fitzgibbon, and Roberto Cipolla. Creatures great and small: Recovering the shape and motion of animals from video. In *Asian Conference on Computer Vision*, pages 3–19. Springer, 2018.
- [7] Piotr Borycki, Yizhe Zhu, Yongqiang Gao, Przemysław Spurek, et al. Smal-pets: Smal based avatars of pets from single image. *arXiv preprint arXiv:2603.17131*, 2026.
- [8] Jinkun Cao, Hongyang Tang, Hao-Shu Fang, Xiaoyong Shen, Cewu Lu, and Yu-Wing Tai. Cross-domain adaptation for animal pose estimation. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 9498–9507, 2019.
- [9] Nicolas Carion, Laura Gustafson, Yuan-Ting Hu, Shoubhik Debnath, Ronghang Hu, Didac Suris, Chaitanya Ryali, Kalyan Vasudev Alwala, Haitham Khedr, Andrew Huang, Jie Lei, Tengyu Ma, Baishan Guo, Arpit Kalla, Markus Marks, Joseph Greer, Meng Wang, Peize Sun, Roman Rädle, Triantafyllos Afouras, Effrosyni Mavroudi, Katherine Xu, Tsung-Han Wu, Yu Zhou, Liliane Momeni, Rishi Hazra, Shuangrui Ding, Sagar Vaze, Francois Porcher, Feng Li, Siyuan Li, Aishwarya Kamath, Ho Kei Cheng, Piotr Dollár, Nikhila Ravi, Kate Saenko, Pengchuan Zhang, and Christoph Feichtenhofer. Sam 3: Segment anything with concepts, 2026.
- [10] Nicolas Carion, Francisco Massa, Gabriel Synnaeve, Nicolas Usunier, Alexander Kirillov, and Sergey Zagoruyko. End-to-end object detection with transformers. In *European conference on computer vision*, pages 213–229. Springer, 2020.
- [11] Thomas J Cashman and Andrew W Fitzgibbon. What shape are dolphins? building 3d morphable models from 2d images. *IEEE transactions on pattern analysis and machine intelligence*, 35(1):232–244, 2012.
- [12] Alexey Dosovitskiy. An image is worth 16x16 words: Transformers for image recognition at scale. *arXiv preprint arXiv:2010.11929*, 2020.
- [13] Mingqi Gao, Yunqi Miao, and Jungong Han. Sam-body4d: Training-free 4d human body mesh recovery from videos. *arXiv preprint arXiv:2512.08406*, 2025.
- [14] Shubham Goel, Angjoo Kanazawa, and Jitendra Malik. Shape and viewpoint without keypoints. In *European Conference on Computer Vision*, pages 88–104. Springer, 2020.
- [15] Shubham Goel, Georgios Pavlakos, Jathushan Rajasegaran, Angjoo Kanazawa, and Jitendra Malik. Humans in 4d: Reconstructing and tracking humans with transformers. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 14783–14794, 2023.
- [16] Tomas Jakab, Ruining Li, Shangzhe Wu, Christian Rupprecht, and Andrea Vedaldi. Farm3d: Learning articulated 3d animals by distilling 2d diffusion. In *2024 International Conference on 3D Vision (3DV)*, pages 852–861. IEEE, 2024.

- [17] Bo Jin, Jin Lyu, Bin Zhang, Tao Yu, Liang An, Yebin Liu, Meili Wang, et al. Monocular mesh recovery and body measurement of female saanen goats. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pages 38670–38678, 2026.
- [18] Angjoo Kanazawa, Michael J Black, David W Jacobs, and Jitendra Malik. End-to-end recovery of human shape and pose. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 7122–7131, 2018.
- [19] Harold W Kuhn. The hungarian method for the assignment problem. *Naval research logistics quarterly*, 2(1-2):83–97, 1955.
- [20] Peter Kulits, Michael J Black, and Silvia Zuffi. Reconstructing animals and the wild. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pages 16565–16577, 2025.
- [21] Ci Li, Nima Ghorbani, Sofia Broomé, Maheen Rashid, Michael J Black, Elin Hernlund, Hedvig Kjellström, and Silvia Zuffi. hsmal: Detailed horse shape and pose reconstruction for motion pattern recognition. *arXiv preprint arXiv:2106.10102*, 2021.
- [22] Feng Li, Hao Zhang, Shilong Liu, Jian Guo, Lionel M Ni, and Lei Zhang. Dn-detr: Accelerate detr training by introducing query denoising. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 13619–13627, 2022.
- [23] Hongjie Li, Heng Yu, Jiaman Li, Hong-Xing Yu, Ehsan Adeli, C Karen Liu, and Jiajun Wu. Anylift: Scaling motion reconstruction from internet videos via 2d diffusion. *arXiv preprint arXiv:2604.17818*, 2026.
- [24] Zizhang Li, Dor Litvak, Ruining Li, Yunzhi Zhang, Tomas Jakab, Christian Rupprecht, Shangzhe Wu, Andrea Vedaldi, and Jiajun Wu. Learning the 3d fauna of the web. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 9752–9762, 2024.
- [25] Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *European conference on computer vision*, pages 740–755. Springer, 2014.
- [26] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. *arXiv preprint arXiv:1711.05101*, 2017.
- [27] Jin Lyu, Liang An, Pujin Cheng, Yebin Liu, and Xiaoying Tang. 4dequene: Disentangling motion and appearance for 4d equine reconstruction from monocular video. *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2026.
- [28] Jin Lyu, Tianyi Zhu, Yi Gu, Li Lin, Pujin Cheng, Yebin Liu, Xiaoying Tang, and Liang An. Animer: Animal pose and shape estimation using family aware transformer. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pages 17486–17496, 2025.
- [29] Xun Long Ng, Kian Eng Ong, Qichen Zheng, Yun Ni, Si Yong Yeo, and Jun Liu. Animal kingdom: A large and diverse dataset for animal behavior understanding. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 19023–19034, 2022.
- [30] Tomasz Niewiadomski, Anastasios Yiannakidis, Hanz Cuevas-Velasquez, Soubhik Sanyal, Michael J Black, Silvia Zuffi, and Peter Kulits. Generative zoo. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 8492–8502, 2025.
- [31] Hamid Rezatofighi, Nathan Tsoi, JunYoung Gwak, Amir Sadeghian, Ian Reid, and Silvio Savarese. Generalized intersection over union: A metric and a loss for bounding box regression. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 658–666, 2019.
- [32] Nadine Rueegg, Silvia Zuffi, Konrad Schindler, and Michael J Black. Barc: Learning to regress 3d dog shape from images by exploiting breed information. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 3876–3884, 2022.

- [33] Chi Su, Xiaoxuan Ma, Jiajun Su, and Yizhou Wang. Sat-hmr: Real-time multi-person 3d mesh estimation via scale-adaptive tokens. In *Proceedings of the Computer Vision and Pattern Recognition Conference (CVPR)*, pages 16796–16806, June 2025.
- [34] Qwen Team. Qwen3 technical report, 2025.
- [35] Yufu Wang, Yu Sun, Priyanka Patel, Kostas Daniilidis, Michael J Black, and Muhammed Kocabas. Prompthmr: Promptable human mesh recovery. In *Proceedings of the computer vision and pattern recognition conference*, pages 1148–1159, 2025.
- [36] Yufu Wang, Ziyun Wang, Lingjie Liu, and Kostas Daniilidis. Tram: Global trajectory and motion of 3d humans from in-the-wild videos. In *European Conference on Computer Vision*, pages 467–487. Springer, 2024.
- [37] Chenfei Wu, Jiahao Li, Jingren Zhou, Junyang Lin, Kaiyuan Gao, Kun Yan, Sheng ming Yin, Shuai Bai, Xiao Xu, Yilei Chen, Yuxiang Chen, Zecheng Tang, Zekai Zhang, Zhengyi Wang, An Yang, Bowen Yu, Chen Cheng, Dayiheng Liu, Deqing Li, Hang Zhang, Hao Meng, Hu Wei, Jingyuan Ni, Kai Chen, Kuan Cao, Liang Peng, Lin Qu, Minggang Wu, Peng Wang, Shuting Yu, Tingkun Wen, Wensen Feng, Xiaoxiao Xu, Yi Wang, Yichang Zhang, Yongqiang Zhu, Yujia Wu, Yuxuan Cai, and Zenan Liu. Qwen-image technical report, 2025.
- [38] Shangzhe Wu, Tomas Jakab, Christian Rupprecht, and Andrea Vedaldi. Dove: Learning deformable 3d objects by watching videos. *International Journal of Computer Vision*, 131(10):2623–2634, 2023.
- [39] Shangzhe Wu, Ruining Li, Tomas Jakab, Christian Rupprecht, and Andrea Vedaldi. Magicpony: Learning articulated 3d animals in the wild. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 8792–8802, 2023.
- [40] Shangzhe Wu, Ameesh Makadia, Jiajun Wu, Noah Snavely, Richard Tucker, and Angjoo Kanazawa. De-rendering the world’s revolutionary artefacts. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 6338–6347, 2021.
- [41] Jiacong Xu, Yi Zhang, Jiawei Peng, Wufei Ma, Artur Jesslen, Pengliang Ji, Qixin Hu, Jiehua Zhang, Qihao Liu, Jiahao Wang, et al. Animal3d: A comprehensive dataset of 3d animal pose and shape. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 9099–9109, 2023.
- [42] Yufei Xu, Jing Zhang, Qiming Zhang, and Dacheng Tao. Vitpose: Simple vision transformer baselines for human pose estimation. *Advances in neural information processing systems*, 35:38571–38584, 2022.
- [43] Gengshan Yang, Deqing Sun, Varun Jampani, Daniel Vlasic, Forrester Cole, Ce Liu, and Deva Ramanan. Viser: Video-specific surface embeddings for articulated 3d shape reconstruction. *Advances in Neural Information Processing Systems*, 34:19326–19338, 2021.
- [44] Xitong Yang, Devansh Kukreja, Don Pinkus, Anushka Sagar, Taosha Fan, Jinhyung Park, Soyong Shin, Jinkun Cao, Jiawei Liu, Nicolas Ugrinovic, Matt Feiszli, Jitendra Malik, Piotr Dollar, and Kris Kitani. Sam 3d body: Robust full-body human mesh recovery, 2026.
- [45] Yuxiang Yang, Yingqi Deng, Yufei Xu, and Jing Zhang. Aptv2: Benchmarking animal pose estimation and tracking with a large-scale dataset and beyond, 2023.
- [46] Yuxiang Yang, Junjie Yang, Yufei Xu, Jing Zhang, Long Lan, and Dacheng Tao. Apt-36k: A large-scale benchmark for animal pose estimation and tracking. *Advances in Neural Information Processing Systems*, 35:17301–17313, 2022.
- [47] Chun-Han Yao, Wei-Chih Hung, Yuanzhen Li, Michael Rubinstein, Ming-Hsuan Yang, and Varun Jampani. Lassie: Learning articulated shapes from sparse image ensemble via 3d part discovery. *Advances in Neural Information Processing Systems*, 35:15296–15308, 2022.
- [48] Hongwen Zhang, Yating Tian, Yuxiang Zhang, Mengcheng Li, Liang An, Zhenan Sun, and Yebin Liu. Pymaf-x: Towards well-aligned full-body model regression from monocular images. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 45(10):12287–12303, 2023.

- [49] Hongwen Zhang, Yating Tian, Xinchu Zhou, Wanli Ouyang, Yebin Liu, Limin Wang, and Zhenan Sun. Pymaf: 3d human pose and shape regression with pyramidal mesh alignment feedback loop. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 11446–11456, 2021.
- [50] Silvia Zuffi and Michael J Black. Awol: Analysis without synthesis using language. In *European Conference on Computer Vision*, pages 1–19. Springer, 2024.
- [51] Silvia Zuffi, Angjoo Kanazawa, and Michael J Black. Lions and tigers and bears: Capturing non-rigid, 3d, articulated shape from images. In *Proceedings of the IEEE conference on Computer Vision and Pattern Recognition*, pages 3955–3963, 2018.
- [52] Silvia Zuffi, Angjoo Kanazawa, David W Jacobs, and Michael J Black. 3d menagerie: Modeling the 3d shape and pose of animals. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 6365–6373, 2017.
- [53] Silvia Zuffi, Ylva Mellbin, Ci Li, Markus Hoeschle, Hedvig Kjellström, Senya Polikovsky, Elin Hernlund, and Michael J Black. Varen: Very accurate and realistic equine network. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 5374–5383, 2024.

A Bipartite Matching Details

We formulate the assignment between the predicted entities and ground-truth targets as a set prediction problem. Let $\hat{\mathcal{Y}} = \{\hat{y}_i\}_{i=1}^N$ denote the set of predictions, where each prediction comprises bounding box coordinates, a confidence score, and 2D keypoint locations. Let $\mathcal{Y} = \{y_j\}_{j=1}^M$ denote the set of ground-truth targets ($M \leq N$). Following DETR [10], the assignment is determined by searching for an optimal injection $\hat{\sigma} : \{1, \dots, M\} \rightarrow \{1, \dots, N\}$ that minimizes the overall bipartite matching cost:

$$\hat{\sigma} = \arg \min_{\sigma} \sum_{i=1}^M \mathcal{C}_{\text{match}}(\hat{y}_{\sigma(i)}, y_i) \quad (7)$$

where $\sigma(i) \in \{1, \dots, N\}$ is the prediction index assigned to the i -th ground-truth target, and $\mathcal{C}_{\text{match}}(\hat{y}_{\sigma(i)}, y_i)$ is the pair-wise assignment cost between prediction $\sigma(i)$ and ground-truth target i . To ensure that the assigned predictions maintain high localization precision and structural pose alignment, our matching cost is a weighted composition of four distinct terms:

$$\mathcal{C}_{\text{match}} = \lambda_{\text{conf}} \mathcal{C}_{\text{conf}} + \lambda_{\text{bbox}} \mathcal{C}_{\text{bbox}} + \lambda_{\text{giou}} \mathcal{C}_{\text{giou}} + \lambda_{\text{kpts}} \mathcal{C}_{\text{kpts}} \quad (8)$$

Specifically, the weighting factors in the matching cost (Eq. 8) are set as follows: $\lambda_{\text{conf}} = 1$, $\lambda_{\text{bbox}} = 1$, $\lambda_{\text{giou}} = 1$, and $\lambda_{\text{kpts}} = 10$. The terms are defined as follows:

- **Bounding Box Cost ($\mathcal{C}_{\text{bbox}}$):** The $\mathcal{L}_1$ distance between the normalized center coordinates, width, and height of the predicted and ground-truth bounding boxes.
- **GIoU Cost ($\mathcal{C}_{\text{giou}}$):** The Generalized Intersection over Union (GIoU) cost [31], which provides a scale-invariant geometric alignment constraint for the bounding boxes.
- **Confidence Cost ($\mathcal{C}_{\text{conf}}$):** A focal-style penalty applied to the predicted confidence score $\hat{c}_i \in [0, 1]$. It is formulated as $\alpha(1 - \hat{c}_i)^\gamma(-\log \hat{c}_i)$ with focusing parameters $\alpha = 0.25$ and $\gamma = 2.0$ following [33], encouraging the model to match ground truth with high-confidence predictions.
- **Keypoint Cost ($\mathcal{C}_{\text{kpts}}$):** A pose-aware geometric constraint computed as the mean $\mathcal{L}_1$ distance between the predicted 2D keypoints and the corresponding ground-truth keypoints, strictly masked to account only for visible joints.

The optimal permutation $\hat{\sigma}$ is computed efficiently using the Hungarian algorithm [19]. Once the optimal assignment is established, the predicted outputs of multiple animals are reordered to perfectly align with the ground-truth sequence, allowing the subsequent calculation of the standard supervisory loss functions.

B More Details about Herd3D

We provide additional qualitative and statistical details of Herd3D in this section. Fig. 8 visualizes representative multi-animal examples with varying group sizes, where the number of visible animals gradually increases from two to eight. These examples illustrate the diversity of species, poses, viewpoints, background environment, and occlusion patterns covered by Herd3D, as well as the corresponding structural cues used for annotation and visualization, including edge maps, depth maps, RGB images, and pose overlays. Table 4 further summarizes the taxonomic composition of the training set by listing the top-10 animal families and representative species within each family. Together, these results show that Herd3D covers both dense multi-animal scenes and broad cross-species diversity, making it suitable for studying scene-centric animal pose understanding.

C Failure Cases in Herd3D Construction.

During the construction of Herd3D, we observed several failure cases when rendering multi-animal scenes with Qwen-ControlNet. These failures are particularly common when animal meshes have large overlapping regions or severe inter-instance occlusions. In such cases, the rendered image may not faithfully preserve the intended 3D geometry and pose. For example, an animal whose head is oriented away from the camera may be incorrectly rendered with a forward-facing face, leading to inconsistent orientation between the mesh and the generated image. Local semantic errors may




Figure 8: Herd3D multi-animal dataset. The images include dogs, horses, antelopes, bears, and cats, with the number of animals per image gradually increasing from 2 to 8 .




Figure 9: Failure cases of data generation.

also occur, where small body parts are misinterpreted, such as ears being rendered as noses or other facial structures. In addition, when two animals are spatially close, the renderer may blend their body regions, causing the torso or limbs of one animal to be partially rendered onto another. These artifacts indicate that dense multi-animal scenes remain challenging for image-conditioned generative rendering, especially under heavy occlusion and close physical interaction. We provide failure cases in Fig. 9.

Table 4: **Top-10 animal families in the Herd3D training image set.** Families are ranked by the number of labelled training images. For each family, we report representative species with their corresponding image counts.

Rank	Family	Images	Representative species
1	Canidae	3,238	Great Dane (105), Curly Coated Retriever (86), Irish Setter (85), Clumber Spaniel (84), Irish Wolfhound (84), Irish Water Spaniel (82), American Cocker Spaniel (73), American Eskimo Dog (71), Cairn Terrier (70), Alaskan Malamute (64)
2	Bovidae	324	Cow (106), Cattle (95), American Bison (21), Antelope (21), Goat (21), Bighorn Sheep (20), Buffalo (20), Impala (20)
3	Equidae	305	Horse (127), Feral Horse (115), Domestic Horse (21), Donkey (21), Common Zebra (21)
4	Mustelidae	187	Honey Badger (73), Eurasian Otter (54), African Clawless Otter (20), American Badger (20), American Mink (20)
5	Rhinocerotidae	157	Black Rhinoceros (157)
6	Ursidae	151	American Black Bear (90), Brown Bear (21), Da Xiong Mao (20), Giant Panda (20)
7	Felidae	143	Bay Lynx (22), European Wildcat (21), Bay Cat (20), Caracal (20), Cat (20), Cheetah (20), Cougar (20)
8	Giraffidae	63	Angolan Giraffe (21), Cape Giraffe (21), Giraffe (21)
9	Herpestidae	42	Gray Meerkat (42)
10	Viverridae	41	Bear Cat (21), African Civet (20)

D More Experiments

Fig. 10 further demonstrates the effect on the number of prompt keypoints. With more keypoints used as prompt information, SAM 3D Animal produces results aligned better with input image.




Figure 10: Ablation study on the number of prompt keypoints.